

#Predict the price of the Uber ride from a given pickup point to the agreed drop-off location. Perform following tasks:

1. Pre-process the dataset.
2. Identify outliers.
3. Check the correlation.
4. Implement linear regression and random forest regression models.
5. Evaluate the models and compare their respective scores like R2, RMSE, etc. Dataset link: <https://www.kaggle.com/datasets/yasserh/uber-fares-dataset>

```
In [1]: #Importing the required libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: #importing the dataset
df = pd.read_csv("uber.csv")
```

1. Pre-process the dataset.

```
In [3]: df.head()
```

```
Out[3]:
```

	Unnamed: 0	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude
0	24238194	2015-05-07 19:52:06.00000003	7.5	2015-05-07 19:52:06 UTC	-73.999817	40.738354
1	27835199	2009-07-17 20:04:56.00000002	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.728225
2	44984355	2009-08-24 21:45:00.000000061	12.9	2009-08-24 21:45:00 UTC	-74.005043	40.740770
3	25894730	2009-06-26 08:22:21.00000001	5.3	2009-06-26 08:22:21 UTC	-73.976124	40.790844
4	17610152	2014-08-28 17:47:00.0000000188	16.0	2014-08-28 17:47:00 UTC	-73.925023	40.744085

```
In [4]: df.info() #To get the required information of the dataset
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            200000 non-null  int64
1   key                   200000 non-null  object
2   fare_amount           200000 non-null  float64
3   pickup_datetime       200000 non-null  object
4   pickup_longitude      200000 non-null  float64
5   pickup_latitude       200000 non-null  float64
6   dropoff_longitude     199999 non-null  float64
7   dropoff_latitude      199999 non-null  float64
8   passenger_count       200000 non-null  int64
dtypes: float64(5), int64(2), object(2)
memory usage: 13.7+ MB
```

In [5]: `df.columns` *#To get number of columns in the dataset*

Out[5]: Index(['Unnamed: 0', 'key', 'fare_amount', 'pickup_datetime', 'pickup_longitude', 'pickup_latitude', 'dropoff_longitude', 'dropoff_latitude', 'passenger_count'], dtype='object')

In [6]: `df = df.drop(['Unnamed: 0', 'key'], axis=1)` *#To drop unnamed column as it isn't relevant*

In [7]: `df.head()`

Out[7]:

	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
0	7.5	2015-05-07 19:52:06 UTC	-73.999817	40.738354	-73.999512	40.738354
1	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.728225	-73.994710	40.728225
2	12.9	2009-08-24 21:45:00 UTC	-74.005043	40.740770	-73.962565	40.740770
3	5.3	2009-06-26 08:22:21 UTC	-73.976124	40.790844	-73.965316	40.790844
4	16.0	2014-08-28 17:47:00 UTC	-73.925023	40.744085	-73.973082	40.744085

In [8]: `df.shape` *#To get the total (Rows,Columns)*

Out[8]: (200000, 7)

In [9]: `df.dtypes` *#To get the type of each column*

Out[9]: fare_amount float64
pickup_datetime object
pickup_longitude float64
pickup_latitude float64
dropoff_longitude float64
dropoff_latitude float64
passenger_count int64
dtype: object

In [10]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fare_amount           200000 non-null float64
1   pickup_datetime       200000 non-null object
2   pickup_longitude      200000 non-null float64
3   pickup_latitude       200000 non-null float64
4   dropoff_longitude     199999 non-null float64
5   dropoff_latitude      199999 non-null float64
6   passenger_count       200000 non-null int64
dtypes: float64(5), int64(1), object(1)
memory usage: 10.7+ MB
```

In [11]: `df.describe()` *#To get statistics of each columns*

Out[11]:

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
count	200000.000000	200000.000000	200000.000000	199999.000000	199999.000000	200000.000000
mean	11.359955	-72.527638	39.935885	-72.525292	39.923890	1.716961
std	9.901776	11.437787	7.720539	13.117408	6.794829	1.468343
min	-52.000000	-1340.648410	-74.015515	-3356.666300	-881.985513	0
25%	6.000000	-73.992065	40.734796	-73.991407	40.733823	1
50%	8.500000	-73.981823	40.752592	-73.980093	40.753042	1
75%	12.500000	-73.967154	40.767158	-73.963658	40.768001	2
max	499.000000	57.418457	1644.421482	1153.572603	872.697628	16

Filling Missing values

In [12]: `df.isnull().sum()`

Out[12]:

```
fare_amount           0
pickup_datetime       0
pickup_longitude      0
pickup_latitude       0
dropoff_longitude     1
dropoff_latitude      1
passenger_count       0
dtype: int64
```

In [13]: `df['dropoff_latitude'].fillna(value=df['dropoff_latitude'].mean(),inplace = True)`
`df['dropoff_longitude'].fillna(value=df['dropoff_longitude'].median(),inplace = True)`

In [14]: `df.isnull().sum()`

```
Out[14]: fare_amount      0
         pickup_datetime  0
         pickup_longitude  0
         pickup_latitude   0
         dropoff_longitude  0
         dropoff_latitude  0
         passenger_count   0
         dtype: int64
```

```
In [15]: df.dtypes
```

```
Out[15]: fare_amount      float64
         pickup_datetime   object
         pickup_longitude  float64
         pickup_latitude   float64
         dropoff_longitude  float64
         dropoff_latitude  float64
         passenger_count   int64
         dtype: object
```

Column pickup_datetime is in wrong format (Object). Convert it to DateTime Format

```
In [16]: df.pickup_datetime = pd.to_datetime(df.pickup_datetime, errors='coerce')
```

```
In [17]: df.dtypes
```

```
Out[17]: fare_amount      float64
         pickup_datetime  datetime64[ns, UTC]
         pickup_longitude  float64
         pickup_latitude   float64
         dropoff_longitude  float64
         dropoff_latitude  float64
         passenger_count   int64
         dtype: object
```

To segregate each time of date and time

```
In [18]: df = df.assign(hour = df.pickup_datetime.dt.hour,
                        day = df.pickup_datetime.dt.day,
                        month = df.pickup_datetime.dt.month,
                        year = df.pickup_datetime.dt.year,
                        dayofweek = df.pickup_datetime.dt.dayofweek)
```

```
In [19]: df.head()
```

Out[19]:

	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_la
0	7.5	2015-05-07 19:52:06+00:00	-73.999817	40.738354	-73.999512	40.7
1	7.7	2009-07-17 20:04:56+00:00	-73.994355	40.728225	-73.994710	40.7
2	12.9	2009-08-24 21:45:00+00:00	-74.005043	40.740770	-73.962565	40.7
3	5.3	2009-06-26 08:22:21+00:00	-73.976124	40.790844	-73.965316	40.8
4	16.0	2014-08-28 17:47:00+00:00	-73.925023	40.744085	-73.973082	40.7

In [20]:

```
# drop the column 'pickup_datetime' using drop()
# 'axis = 1' drops the specified column

df = df.drop('pickup_datetime',axis=1)
```

In [21]:

```
df.head()
```

Out[21]:

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger
0	7.5	-73.999817	40.738354	-73.999512	40.723217	
1	7.7	-73.994355	40.728225	-73.994710	40.750325	
2	12.9	-74.005043	40.740770	-73.962565	40.772647	
3	5.3	-73.976124	40.790844	-73.965316	40.803349	
4	16.0	-73.925023	40.744085	-73.973082	40.761247	

In [22]:

```
df.dtypes
```

Out[22]:

```
fare_amount          float64
pickup_longitude      float64
pickup_latitude       float64
dropoff_longitude     float64
dropoff_latitude      float64
passenger_count       int64
hour                  int64
day                   int64
month                 int64
year                  int64
dayofweek             int64
dtype: object
```

Checking outliers and filling them


In [23]:

```
df.plot(kind = "box",subplots = True,layout = (7,2),figsize=(15,20)) #BoxPlot to ch
```

```

Out[23]: fare_amount      AxesSubplot(0.125,0.787927;0.352273x0.0920732)
pickup_longitude      AxesSubplot(0.547727,0.787927;0.352273x0.0920732)
pickup_latitude      AxesSubplot(0.125,0.677439;0.352273x0.0920732)
dropoff_longitude      AxesSubplot(0.547727,0.677439;0.352273x0.0920732)
dropoff_latitude      AxesSubplot(0.125,0.566951;0.352273x0.0920732)
passenger_count      AxesSubplot(0.547727,0.566951;0.352273x0.0920732)
hour      AxesSubplot(0.125,0.456463;0.352273x0.0920732)
day      AxesSubplot(0.547727,0.456463;0.352273x0.0920732)
month      AxesSubplot(0.125,0.345976;0.352273x0.0920732)
year      AxesSubplot(0.547727,0.345976;0.352273x0.0920732)
dayofweek      AxesSubplot(0.125,0.235488;0.352273x0.0920732)
dtype: object

```



```

In [24]: #Using the InterQuartile Range to fill the values
def remove_outlier(df1 , col):
    Q1 = df1[col].quantile(0.25)
    Q3 = df1[col].quantile(0.75)
    IQR = Q3 - Q1
    lower_whisker = Q1-1.5*IQR
    upper_whisker = Q3+1.5*IQR
    df[col] = np.clip(df1[col] , lower_whisker , upper_whisker)
    return df1

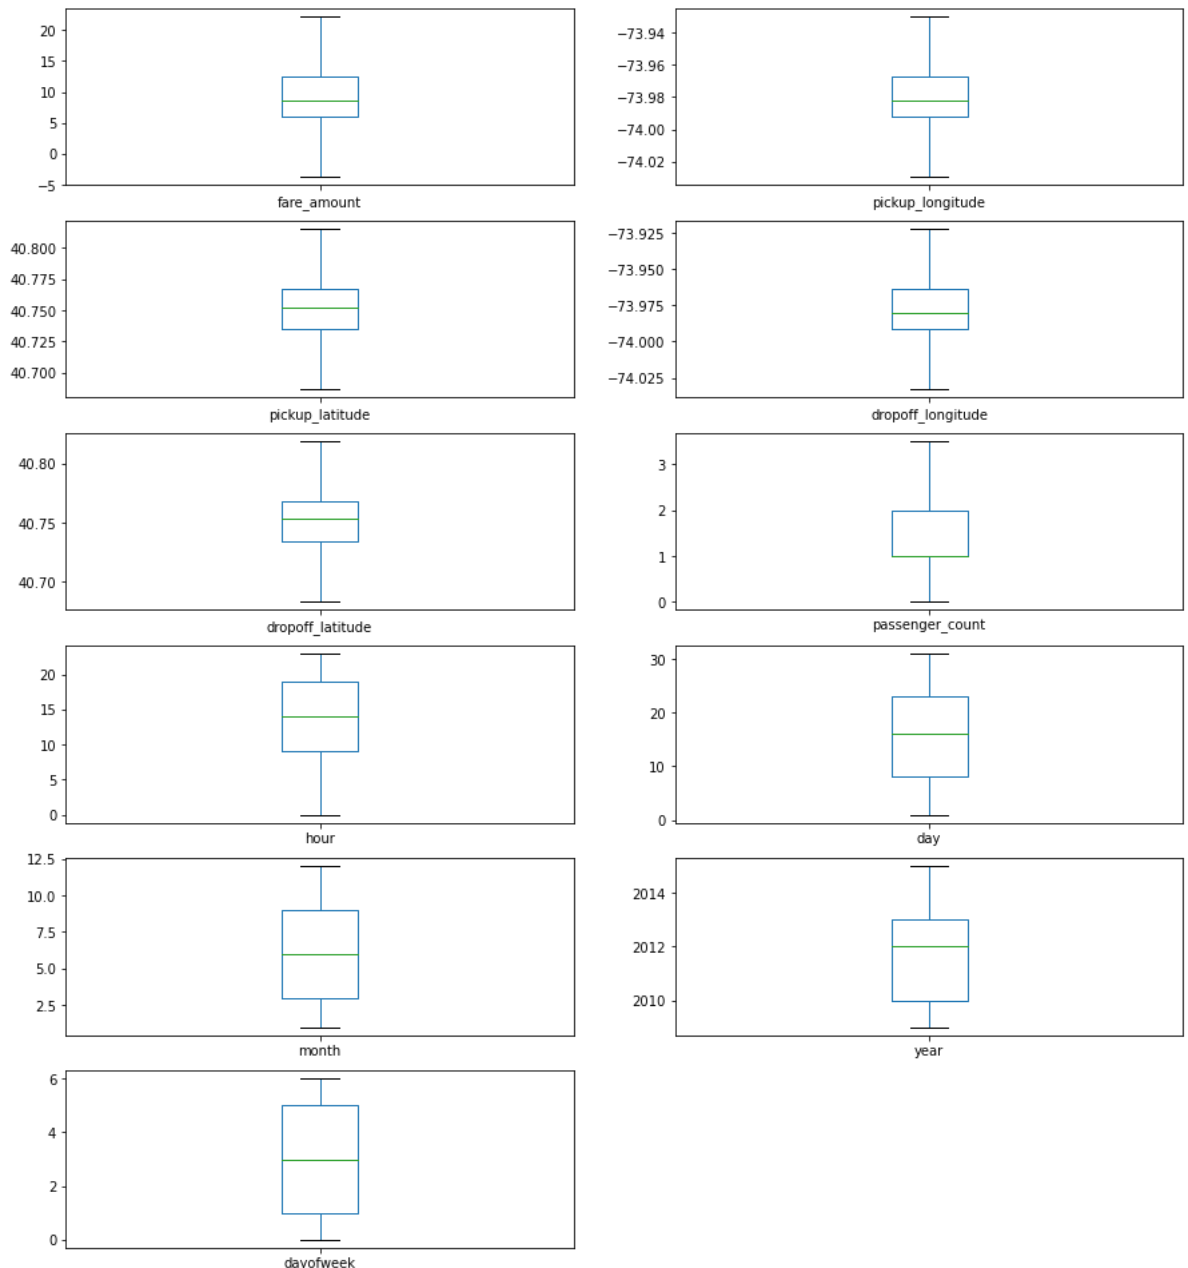
```

```
def treat_outliers_all(df1 , col_list):
    for c in col_list:
        df1 = remove_outlier(df , c)
    return df1
```

```
In [25]: df = treat_outliers_all(df , df.iloc[:, 0::])
```

```
In [26]: df.plot(kind = "box",subplots = True,layout = (7,2),figsize=(15,20)) #BoxPlot shows
```

```
Out[26]: fare_amount      AxesSubplot(0.125,0.787927;0.352273x0.0920732)
pickup_longitude      AxesSubplot(0.547727,0.787927;0.352273x0.0920732)
pickup_latitude      AxesSubplot(0.125,0.677439;0.352273x0.0920732)
dropoff_longitude      AxesSubplot(0.547727,0.677439;0.352273x0.0920732)
dropoff_latitude      AxesSubplot(0.125,0.566951;0.352273x0.0920732)
passenger_count      AxesSubplot(0.547727,0.566951;0.352273x0.0920732)
hour      AxesSubplot(0.125,0.456463;0.352273x0.0920732)
day      AxesSubplot(0.547727,0.456463;0.352273x0.0920732)
month      AxesSubplot(0.125,0.345976;0.352273x0.0920732)
year      AxesSubplot(0.547727,0.345976;0.352273x0.0920732)
dayofweek      AxesSubplot(0.125,0.235488;0.352273x0.0920732)
dtype: object
```



```
In [27]: #pip install haversine
```

```
import haversine as hs #Calculate the distance using Haversine to calculate the di
travel_dist = []
for pos in range(len(df['pickup_longitude'])):
    long1,lati1,long2,lati2 = [df['pickup_longitude'][pos],df['pickup_latitude'
    loc1=(lati1,long1)
    loc2=(lati2,long2)
    c = hs.haversine(loc1,loc2)
    travel_dist.append(c)

print(travel_dist)
df['dist_travel_km'] = travel_dist
df.head()
```

IOPub data rate exceeded.
The notebook server will temporarily stop sending output
to the client in order to avoid crashing it.
To change this limit, set the config variable
`--NotebookApp.iopub_data_rate_limit`.

Current values:
NotebookApp.iopub_data_rate_limit=1000000.0 (bytes/sec)
NotebookApp.rate_limit_window=3.0 (secs)

Out[27]:

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger
0	7.5	-73.999817	40.738354	-73.999512	40.723217	
1	7.7	-73.994355	40.728225	-73.994710	40.750325	
2	12.9	-74.005043	40.740770	-73.962565	40.772647	
3	5.3	-73.976124	40.790844	-73.965316	40.803349	
4	16.0	-73.929786	40.744085	-73.973082	40.761247	

In [28]: *#Uber doesn't travel over 130 kms so minimize the distance*
df= df.loc[(df.dist_travel_km >= 1) | (df.dist_travel_km <= 130)]
print("Remaining observastions in the dataset:", df.shape)

Remaining observastions in the dataset: (200000, 12)

In [29]: *#Finding inccorect Latitude (Less than or greater than 90) and Longitude (greater t*
incorrect_coordinates = df.loc[(df.pickup_latitude > 90) |(df.pickup_latitude < -90
(df.dropoff_latitude > 90) |(df.dropoff_latitude
(df.pickup_longitude > 180) |(df.pickup_longitu
(df.dropoff_longitude > 90) |(df.dropoff_longitu
]

In [30]: df.drop(incorrect_coordinates, inplace = True, errors = 'ignore')

In [31]: df.head()


```
Out[31]:
```


	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger
0	7.5	-73.999817	40.738354	-73.999512	40.723217	
1	7.7	-73.994355	40.728225	-73.994710	40.750325	
2	12.9	-74.005043	40.740770	-73.962565	40.772647	
3	5.3	-73.976124	40.790844	-73.965316	40.803349	
4	16.0	-73.929786	40.744085	-73.973082	40.761247	

```
In [32]: df.isnull().sum()
```

```
Out[32]: fare_amount      0
pickup_longitude    0
pickup_latitude     0
dropoff_longitude   0
dropoff_latitude    0
passenger_count     0
hour                0
day                 0
month               0
year                0
dayofweek           0
dist_travel_km      0
dtype: int64
```

```
In [33]: sns.heatmap(df.isnull()) #Free for null values
```

```
Out[33]: <AxesSubplot:>
```



```
In [34]: corr = df.corr() #Function to find the correlation
```

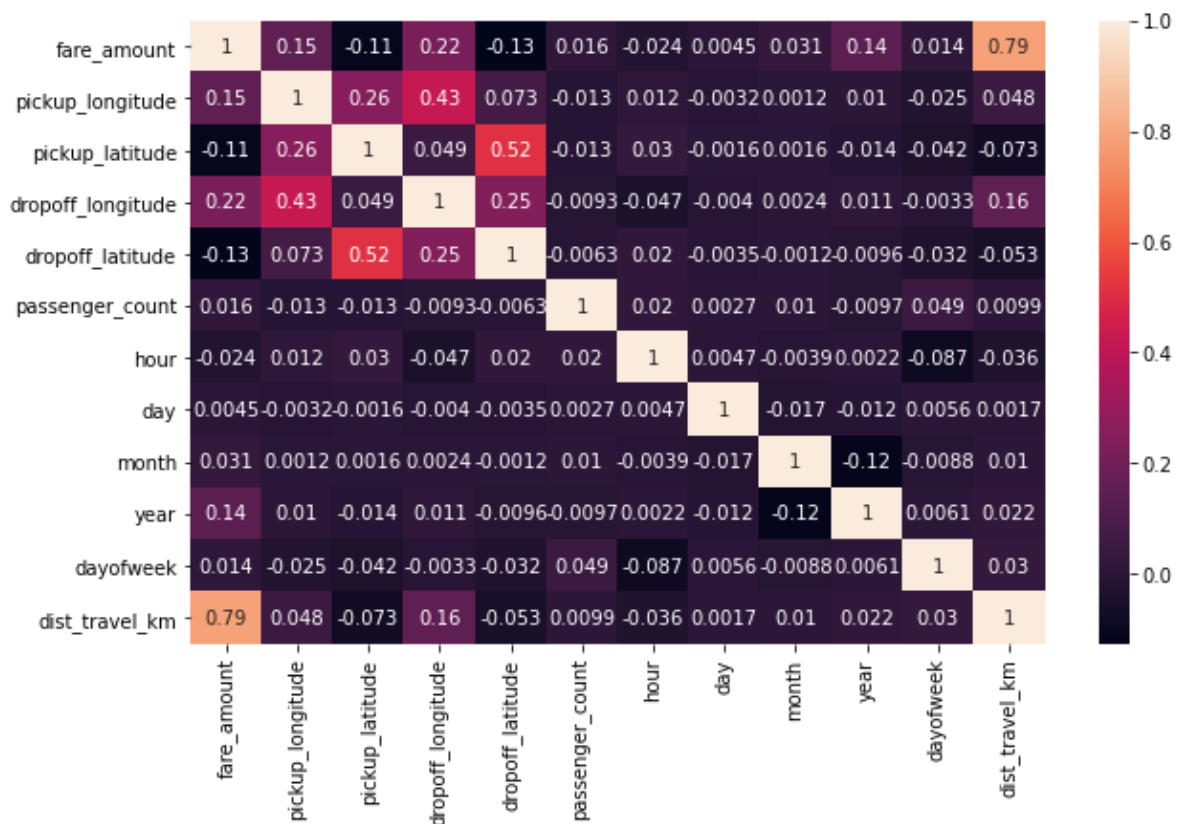
```
In [35]: corr
```

Out[35]:

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
fare_amount	1.000000	0.154069	-0.110842	0.218675	-0.12
pickup_longitude	0.154069	1.000000	0.259497	0.425619	0.07
pickup_latitude	-0.110842	0.259497	1.000000	0.048889	0.51
dropoff_longitude	0.218675	0.425619	0.048889	1.000000	0.24
dropoff_latitude	-0.125898	0.073290	0.515714	0.245667	1.00
passenger_count	0.015778	-0.013213	-0.012889	-0.009303	-0.00
hour	-0.023623	0.011579	0.029681	-0.046558	0.01
day	0.004534	-0.003204	-0.001553	-0.004007	-0.00
month	0.030817	0.001169	0.001562	0.002391	-0.00
year	0.141277	0.010198	-0.014243	0.011346	-0.00
dayofweek	0.013652	-0.024652	-0.042310	-0.003336	-0.03
dist_travel_km	0.786385	0.048446	-0.073362	0.155191	-0.05

In [36]: `fig,axis = plt.subplots(figsize = (10,6))`
`sns.heatmap(df.corr(),annot = True)` *#Correlation Heatmap (Light values means highly*

Out[36]: <AxesSubplot:>



Dividing the dataset into feature and target values

In [182]: `x = df[['pickup_longitude','pickup_latitude','dropoff_longitude','dropoff_latitude']]`

In [183]: `y = df['fare_amount']`

Dividing the dataset into training and testing dataset

```
In [184... from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(x,y,test_size = 0.33)
```

Linear Regression

```
In [185... from sklearn.linear_model import LinearRegression
regression = LinearRegression()
```

```
In [186... regression.fit(X_train,y_train)
```

Out[186]: LinearRegression()

```
In [80]: regression.intercept_ #To find the linear intercept
```

Out[80]: 2640.1356169149753

```
In [187... regression.coef_ #To find the linear coefficient
```

Out[187]: array([2.54805415e+01, -7.18365435e+00, 1.96232986e+01, -1.79401980e+01,
 5.48472723e-02, 5.32910041e-03, 4.05930990e-03, 5.74261856e-02,
 3.66574831e-01, -3.03753790e-02, 1.84233728e+00])

```
In [188... prediction = regression.predict(X_test) #To predict the target values
```

```
In [189... print(prediction)
```

```
[ 5.47848314 10.11016249 12.19490542 ...  7.11952609 20.2482979
 8.82791961]
```

```
In [190... y_test
```

Out[190]: 155740 4.90
47070 10.00
116192 14.50
164589 6.50
154309 11.30
...
76552 7.70
27926 10.90
38972 6.50
120341 22.25
178449 8.10
Name: fare_amount, Length: 66000, dtype: float64

Metrics Evaluation using R2, Mean Squared Error, Root Mean Squared Error

```
In [191... from sklearn.metrics import r2_score
```

```
In [192... r2_score(y_test,prediction)
```

Out[192]: 0.6651880468683617

```
In [193... from sklearn.metrics import mean_squared_error
```

```
In [194...] MSE = mean_squared_error(y_test,prediction)
```

```
In [195...] MSE
```

```
Out[195]: 9.961516917717704
```

```
In [196...] RMSE = np.sqrt(MSE)
```

```
In [197...] RMSE
```

```
Out[197]: 3.156187085348032
```

Random Forest Regression

```
In [198...] from sklearn.ensemble import RandomForestRegressor
```

```
In [199...] rf = RandomForestRegressor(n_estimators=100) #Here n_estimators means number of trees
```

```
In [200...] rf.fit(X_train,y_train)
```

```
Out[200]: RandomForestRegressor()
```

```
In [201...] y_pred = rf.predict(X_test)
```

```
In [202...] y_pred
```

```
Out[202]: array([ 5.714 , 10.285 , 12.68  , ...,  6.338 , 19.4685,  7.712 ])
```

Metrics evaluation for Random Forest

```
In [210...] R2_Random = r2_score(y_test,y_pred)
```

```
In [211...] R2_Random
```

```
Out[211]: 0.7948374920410631
```

```
In [205...] MSE_Random = mean_squared_error(y_test,y_pred)
```

```
In [206...] MSE_Random
```

```
Out[206]: 6.104112397417331
```

```
In [207...] RMSE_Random = np.sqrt(MSE_Random)
```

```
In [208...] RMSE_Random
```

```
Out[208]: 2.4706501972997574
```

Assignment 2

2. Classify the email using the binary classification method. Email Spam detection has two states: a) Normal State – Not Spam, b) Abnormal State – Spam. Use K-Nearest Neighbors and Support Vector Machine for classification. Analyze their performance.

Dataset link: The emails.csv dataset on the Kaggle

<https://www.kaggle.com/datasets/balaka18/email-spam-classification-dataset-csv>

```
In [19]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn import metrics
```

```
In [20]: df=pd.read_csv('emails.csv')
```

```
In [21]: df.head()
```

```
Out[21]:
```

	Email No.	the	to	ect	and	for	of	a	you	hou	...	connevey	jay	valued	lay	infrastruct
0	Email 1	0	0	1	0	0	0	2	0	0	...	0	0	0	0	
1	Email 2	8	13	24	6	6	2	102	1	27	...	0	0	0	0	
2	Email 3	0	0	1	0	0	0	8	0	0	...	0	0	0	0	
3	Email 4	0	5	22	0	5	1	51	2	10	...	0	0	0	0	
4	Email 5	7	6	17	1	5	2	57	0	9	...	0	0	0	0	

5 rows × 3002 columns

```
In [22]: df.columns
```

```
Out[22]: Index(['Email No.', 'the', 'to', 'ect', 'and', 'for', 'of', 'a', 'you', 'hou',
...,
'connevey', 'jay', 'valued', 'lay', 'infrastructure', 'military',
'allowing', 'ff', 'dry', 'Prediction'],
dtype='object', length=3002)
```

```
In [23]: df.isnull().sum()
```

```
Out[23]: Email No.      0
         the           0
         to           0
         ect          0
         and          0
         ..
         military     0
         allowing     0
         ff           0
         dry          0
         Prediction   0
         Length: 3002, dtype: int64
```

```
In [24]: df.dropna(inplace = True)
```

```
In [25]: df.drop(['Email No.'],axis=1,inplace=True)
X = df.drop(['Prediction'],axis = 1)
y = df['Prediction']
```

```
In [26]: from sklearn.preprocessing import scale
X = scale(X)
# split into train and test
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_s
```

KNN classifier

```
In [35]: from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=7)

knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)
```

```
In [36]: print("Prediction",y_pred)

Prediction [0 0 1 ... 1 1 1]
```

```
In [37]: print("KNN accuracy = ",metrics.accuracy_score(y_test,y_pred))

KNN accuracy = 0.8009020618556701
```

```
In [39]: print("Confusion matrix",metrics.confusion_matrix(y_test,y_pred))

Confusion matrix [[804 293]
 [ 16 439]]
```

SVM classifier

```
In [27]: # cost C = 1
model = SVC(C = 1)

# fit
model.fit(X_train, y_train)

# predict
y_pred = model.predict(X_test)
```

```
In [28]: metrics.confusion_matrix(y_true=y_test, y_pred=y_pred)
```

```
Out[28]: array([[1091,    6],  
               [  90,  365]])
```

```
In [29]: print("SVM accuracy = ",metrics.accuracy_score(y_test,y_pred))
```

```
SVM accuracy = 0.9381443298969072
```

Given a bank customer, build a neural network-based classifier that can determine whether they will leave or not in the next 6 months.

Dataset Description: The case study is from an open-source dataset from Kaggle. The dataset contains 10,000 sample points with 14 distinct features such as CustomerId, CreditScore, Geography, Gender, Age, Tenure, Balance, etc. Link to the Kaggle project: <https://www.kaggle.com/barelydedicated/bank-customer-churn-modeling> Perform following steps:

1. Read the dataset.
2. Distinguish the feature and target set and divide the data set into training and test sets.
3. Normalize the train and test data.
4. Initialize and build the model. Identify the points of improvement and implement the same.
5. Print the accuracy score and confusion matrix.

```
In [46]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt #Importing the libraries
```

```
In [47]: df = pd.read_csv("Churn_Modelling.csv")
```

Preprocessing.

```
In [48]: df.head()
```

```
Out[48]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance
0	1	15634602	Hargrave	619	France	Female	42	2	0.00
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86
2	3	15619304	Onio	502	France	Female	42	8	159660.80
3	4	15701354	Boni	699	France	Female	39	1	0.00
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82

```
In [49]: df.shape
```

```
Out[49]: (10000, 14)
```

```
In [50]: df.describe()
```


Out[50]:

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	Num
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	1
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288	
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405202	
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000	
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000	
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.090000	

In [51]:

df.isnull()

Out[51]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balan
0	False	False	False	False	False	False	False	False	Fal
1	False	False	False	False	False	False	False	False	Fal
2	False	False	False	False	False	False	False	False	Fal
3	False	False	False	False	False	False	False	False	Fal
4	False	False	False	False	False	False	False	False	Fal
...
9995	False	False	False	False	False	False	False	False	Fal
9996	False	False	False	False	False	False	False	False	Fal
9997	False	False	False	False	False	False	False	False	Fal
9998	False	False	False	False	False	False	False	False	Fal
9999	False	False	False	False	False	False	False	False	Fal

10000 rows × 14 columns

In [52]:

df.isnull().sum()

Out[52]:

RowNumber0
CustomerId0
Surname0
CreditScore0
Geography0
Gender0
Age0
Tenure0
Balance0
NumOfProducts0
HasCrCard0
IsActiveMember0
EstimatedSalary0
Exited0
dtype: int64

In [53]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   RowNumber             10000 non-null  int64
 1   CustomerId            10000 non-null  int64
 2   Surname               10000 non-null  object
 3   CreditScore           10000 non-null  int64
 4   Geography             10000 non-null  object
 5   Gender               10000 non-null  object
 6   Age                   10000 non-null  int64
 7   Tenure                10000 non-null  int64
 8   Balance               10000 non-null  float64
 9   NumOfProducts         10000 non-null  int64
10   HasCrCard             10000 non-null  int64
11   IsActiveMember        10000 non-null  int64
12   EstimatedSalary       10000 non-null  float64
13   Exited                10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

In [54]: `df.dtypes`

```
Out[54]: RowNumber          int64
CustomerId          int64
Surname             object
CreditScore         int64
Geography           object
Gender              object
Age                 int64
Tenure              int64
Balance             float64
NumOfProducts       int64
HasCrCard           int64
IsActiveMember      int64
EstimatedSalary     float64
Exited              int64
dtype: object
```

In [55]: `df.columns`

```
Out[55]: Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
                'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
                'IsActiveMember', 'EstimatedSalary', 'Exited'],
               dtype='object')
```

In [56]: `df = df.drop(['RowNumber', 'Surname', 'CustomerId'], axis=1) #Dropping the unneces`

In [57]: `df.head()`

Out[57]:


	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActive
0	619	France	Female	42	2	0.00	1	1	
1	608	Spain	Female	41	1	83807.86	1	0	
2	502	France	Female	42	8	159660.80	3	1	
3	699	France	Female	39	1	0.00	2	0	
4	850	Spain	Female	43	2	125510.82	1	1	

Visualization

```
In [101...] def visualization(x, y, xlabel):
    plt.figure(figsize=(10,5))
    plt.hist([x, y], color=['red', 'green'], label = ['exit', 'not_exit'])
    plt.xlabel(xlabel, fontsize=20)
    plt.ylabel("No. of customers", fontsize=20)
    plt.legend()
```


```
In [102...] df_churn_exited = df[df['Exited']==1]['Tenure']
df_churn_not_exited = df[df['Exited']==0]['Tenure']
```

```
In [103...] visualization(df_churn_exited, df_churn_not_exited, "Tenure")
```



```
In [105...] df_churn_exited2 = df[df['Exited']==1]['Age']
df_churn_not_exited2 = df[df['Exited']==0]['Age']
```

```
In [106...] visualization(df_churn_exited2, df_churn_not_exited2, "Age")
```



Converting the Categorical Variables

```
In [59]: X = df[['CreditScore', 'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
states = pd.get_dummies(df['Geography'], drop_first = True)
gender = pd.get_dummies(df['Gender'], drop_first = True)
```

```
In [61]: df = pd.concat([df, gender, states], axis = 1)
```

Splitting the training and testing Dataset

```
In [62]: df.head()
```

```
Out[62]:
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActive
0	619	France	Female	42	2	0.00	1	1	
1	608	Spain	Female	41	1	83807.86	1	0	
2	502	France	Female	42	8	159660.80	3	1	
3	699	France	Female	39	1	0.00	2	0	
4	850	Spain	Female	43	2	125510.82	1	1	

```
In [63]: X = df[['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard', 'IsActive']
```

```
In [64]: y = df['Exited']
```

```
In [65]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.30)
```

Normalizing the values with mean as 0 and Standard Deviation as 1

```
In [66]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
```

```
In [67]: X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

```
In [68]: X_train
```

```
Out[68]: array([[ 4.56838557e-01, -9.45594735e-01,  1.58341939e-03, ...,
                  9.13181783e-01, -5.81969145e-01, -5.73611200e-01],
                [-2.07591864e-02, -2.77416637e-01,  3.47956411e-01, ...,
                 -1.09507222e+00, -5.81969145e-01,  1.74334114e+00],
                [-1.66115021e-01,  1.82257167e+00, -1.38390855e+00, ...,
                 -1.09507222e+00, -5.81969145e-01, -5.73611200e-01],
                ...,
                [-3.63383654e-01, -4.68324665e-01,  1.73344838e+00, ...,
                  9.13181783e-01, -5.81969145e-01, -5.73611200e-01],
                [ 4.67221117e-01, -1.42286480e+00,  1.38707539e+00, ...,
                  9.13181783e-01, -5.81969145e-01,  1.74334114e+00],
                [-8.82511636e-01,  2.95307447e-01, -6.91162564e-01, ...,
                  9.13181783e-01, -5.81969145e-01, -5.73611200e-01]])
```

```
In [69]: X_test
```

```
Out[69]: array([[ 3.63395520e-01,  1.99853433e-01,  1.58341939e-03, ...,
                  9.13181783e-01, -5.81969145e-01, -5.73611200e-01],
                [-4.15243057e-02,  4.86215475e-01,  1.58341939e-03, ...,
                 -1.09507222e+00, -5.81969145e-01,  1.74334114e+00],
                [-1.87923736e+00, -3.72870651e-01, -1.38390855e+00, ...,
                  9.13181783e-01, -5.81969145e-01, -5.73611200e-01],
                ...,
                [-6.02182526e-01, -5.63778679e-01, -1.73028154e+00, ...,
                 -1.09507222e+00, -5.81969145e-01, -5.73611200e-01],
                [ 1.51585964e+00, -6.59232693e-01,  1.73344838e+00, ...,
                  9.13181783e-01, -5.81969145e-01, -5.73611200e-01],
                [-5.19122049e-01,  1.04399419e-01,  1.73344838e+00, ...,
                  9.13181783e-01, -5.81969145e-01, -5.73611200e-01]])
```

Building the Classifier Model using Keras

```
In [70]: import keras #Keras is the wrapper on the top of tensorflow
#Can use Tensorflow as well but won't be able to understand the errors initially.
```

```
In [71]: from keras.models import Sequential #To create sequential neural network
from keras.layers import Dense #To create hidden layers
```

```
In [72]: classifier = Sequential()
```

```
In [74]: #To add the layers
#Dense helps to construct the neurons
#Input Dimension means we have 11 features
#Units is to create the hidden layers
#Uniform helps to distribute the weight uniformly
classifier.add(Dense(activation = "relu",input_dim = 11,units = 6,kernel_initializer = "uniform"))
```

```
In [75]: classifier.add(Dense(activation = "relu",units = 6,kernel_initializer = "uniform"))
```

```
In [76]: classifier.add(Dense(activation = "sigmoid",units = 1,kernel_initializer = "uniform
```

```
In [77]: classifier.compile(optimizer="adam",loss = 'binary_crossentropy',metrics = ['accura
```

```
In [79]: classifier.summary() #3 Layers created. 6 neurons in 1st,6neurons in 2nd Layer and
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
dense_3 (Dense)	(None, 6)	72
=====		
dense_4 (Dense)	(None, 6)	42
=====		
dense_5 (Dense)	(None, 1)	7
=====		
Total params: 121		
Trainable params: 121		
Non-trainable params: 0		
=====		

```
In [89]: classifier.fit(X_train,y_train,batch_size=10,epochs=50) #Fitting the ANN to trainir
```

```
Epoch 1/50
700/700 [=====] - 0s 674us/step - loss: 0.4293 - accurac
y: 0.7947
Epoch 2/50
700/700 [=====] - 0s 647us/step - loss: 0.4239 - accurac
y: 0.7947
Epoch 3/50
700/700 [=====] - 0s 657us/step - loss: 0.4203 - accurac
y: 0.8067
Epoch 4/50
700/700 [=====] - 0s 664us/step - loss: 0.4167 - accurac
y: 0.8260
Epoch 5/50
700/700 [=====] - 0s 674us/step - loss: 0.4153 - accurac
y: 0.8287
Epoch 6/50
700/700 [=====] - 0s 653us/step - loss: 0.4137 - accurac
y: 0.8310
Epoch 7/50
700/700 [=====] - 0s 658us/step - loss: 0.4125 - accurac
y: 0.8317
Epoch 8/50
700/700 [=====] - 1s 842us/step - loss: 0.4116 - accurac
y: 0.8306
Epoch 9/50
700/700 [=====] - 0s 671us/step - loss: 0.4103 - accurac
y: 0.8331
Epoch 10/50
700/700 [=====] - 0s 682us/step - loss: 0.4100 - accurac
y: 0.8326
Epoch 11/50
700/700 [=====] - 0s 690us/step - loss: 0.4093 - accurac
y: 0.8337
Epoch 12/50
700/700 [=====] - 0s 688us/step - loss: 0.4087 - accurac
y: 0.8339
Epoch 13/50
700/700 [=====] - 0s 675us/step - loss: 0.4081 - accurac
y: 0.8341
Epoch 14/50
700/700 [=====] - 1s 722us/step - loss: 0.4071 - accurac
y: 0.8331
Epoch 15/50
700/700 [=====] - 1s 811us/step - loss: 0.4065 - accurac
y: 0.8341
Epoch 16/50
700/700 [=====] - 0s 711us/step - loss: 0.4056 - accurac
y: 0.8356
Epoch 17/50
700/700 [=====] - 0s 702us/step - loss: 0.4046 - accurac
y: 0.8366
Epoch 18/50
700/700 [=====] - 0s 688us/step - loss: 0.4035 - accurac
y: 0.8343
Epoch 19/50
700/700 [=====] - 1s 715us/step - loss: 0.4024 - accurac
y: 0.8363
Epoch 20/50
700/700 [=====] - 0s 714us/step - loss: 0.4020 - accurac
y: 0.8337
Epoch 21/50
```

```
700/700 [=====] - 0s 705us/step - loss: 0.4010 - accurac
y: 0.8374
Epoch 22/50
700/700 [=====] - 1s 720us/step - loss: 0.4003 - accurac
y: 0.8370
Epoch 23/50
700/700 [=====] - 0s 692us/step - loss: 0.3993 - accurac
y: 0.8374
Epoch 24/50
700/700 [=====] - 0s 709us/step - loss: 0.3990 - accurac
y: 0.8356
Epoch 25/50
700/700 [=====] - 1s 871us/step - loss: 0.3984 - accurac
y: 0.8366
Epoch 26/50
700/700 [=====] - 1s 719us/step - loss: 0.3984 - accurac
y: 0.8367
Epoch 27/50
700/700 [=====] - 1s 719us/step - loss: 0.3980 - accurac
y: 0.8366
Epoch 28/50
700/700 [=====] - 0s 695us/step - loss: 0.3981 - accurac
y: 0.8366
Epoch 29/50
700/700 [=====] - 0s 667us/step - loss: 0.3976 - accurac
y: 0.8374
Epoch 30/50
700/700 [=====] - 0s 669us/step - loss: 0.3972 - accurac
y: 0.8373
Epoch 31/50
700/700 [=====] - 0s 670us/step - loss: 0.3970 - accurac
y: 0.8370
Epoch 32/50
700/700 [=====] - 1s 720us/step - loss: 0.3972 - accurac
y: 0.8376
Epoch 33/50
700/700 [=====] - 0s 675us/step - loss: 0.3965 - accurac
y: 0.8367
Epoch 34/50
700/700 [=====] - 0s 680us/step - loss: 0.3961 - accurac
y: 0.8364
Epoch 35/50
700/700 [=====] - 0s 685us/step - loss: 0.3962 - accurac
y: 0.8379
Epoch 36/50
700/700 [=====] - 1s 771us/step - loss: 0.3960 - accurac
y: 0.8370
Epoch 37/50
700/700 [=====] - 1s 1ms/step - loss: 0.3963 - accuracy:
0.8366
Epoch 38/50
700/700 [=====] - 1s 764us/step - loss: 0.3962 - accurac
y: 0.8373
Epoch 39/50
700/700 [=====] - 1s 823us/step - loss: 0.3950 - accurac
y: 0.8384
Epoch 40/50
700/700 [=====] - 1s 759us/step - loss: 0.3956 - accurac
y: 0.8361
Epoch 41/50
700/700 [=====] - 1s 773us/step - loss: 0.3949 - accurac
```



```

y: 0.8366
Epoch 42/50
700/700 [=====] - 0s 695us/step - loss: 0.3953 - accurac
y: 0.8369
Epoch 43/50
700/700 [=====] - 0s 701us/step - loss: 0.3952 - accurac
y: 0.8369
Epoch 44/50
700/700 [=====] - 0s 707us/step - loss: 0.3952 - accurac
y: 0.8366
Epoch 45/50
700/700 [=====] - 0s 680us/step - loss: 0.3955 - accurac
y: 0.8376
Epoch 46/50
700/700 [=====] - 0s 665us/step - loss: 0.3947 - accurac
y: 0.8373
Epoch 47/50
700/700 [=====] - 0s 708us/step - loss: 0.3947 - accurac
y: 0.8371
Epoch 48/50
700/700 [=====] - 0s 681us/step - loss: 0.3944 - accurac
y: 0.8371
Epoch 49/50
700/700 [=====] - 0s 678us/step - loss: 0.3947 - accurac
y: 0.8383
Epoch 50/50
700/700 [=====] - 1s 869us/step - loss: 0.3944 - accurac
y: 0.8370

```

Out[89]: <tensorflow.python.keras.callbacks.History at 0x1fb1eb93df0>

```
In [90]: y_pred = classifier.predict(X_test)
y_pred = (y_pred > 0.5) #Predicting the result
```

```
In [97]: from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
```

```
In [92]: cm = confusion_matrix(y_test, y_pred)
```

```
In [93]: cm
```

```
Out[93]: array([[2328,   72],
               [ 425,  175]], dtype=int64)
```


```
In [94]: accuracy = accuracy_score(y_test, y_pred)
```

```
In [95]: accuracy
```

Out[95]: 0.8343333333333334

```
In [98]: plt.figure(figsize = (10,7))
sns.heatmap(cm, annot = True)
plt.xlabel('Predicted')
plt.ylabel('Truth')
```

Out[98]: Text(69.0, 0.5, 'Truth')



```
In [100... print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.85	0.97	0.90	2400
1	0.71	0.29	0.41	600
accuracy			0.83	3000
macro avg	0.78	0.63	0.66	3000
weighted avg	0.82	0.83	0.81	3000

```
In [ ]:
```

Assignment 5

KNN algorithm on diabetes dataset

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn import metrics
```

```
In [2]: df=pd.read_csv('diabetes.csv')
```

```
In [3]: df.columns
```

```
Out[3]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
              'BMI', 'Pedigree', 'Age', 'Outcome'],
              dtype='object')
```

Check for null values. If present remove null values from the dataset

```
In [4]: df.isnull().sum()
```

```
Out[4]: Pregnancies      0
Glucose      0
BloodPressure  0
SkinThickness  0
Insulin      0
BMI          0
Pedigree     0
Age          0
Outcome      0
dtype: int64
```

```
In [ ]:
```

Outcome is the label/target, other columns are features

```
In [7]: X = df.drop('Outcome',axis = 1)
y = df['Outcome']
```

```
In [8]: from sklearn.preprocessing import scale
X = scale(X)
# split into train and test
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_s
```

```
In [9]: from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=7)

knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)
```

```
In [17]: print("Confusion matrix: ")
cs = metrics.confusion_matrix(y_test,y_pred)
print(cs)
```

```
Confusion matrix:
[[123  28]
 [ 37  43]]
```

```
In [12]: print("Accuracy ",metrics.accuracy_score(y_test,y_pred))
```

```
Accuracy  0.7186147186147186
```

Classification error rate: proportion of instances misclassified over the whole set of instances. Error rate is calculated as the total number of two incorrect predictions (FN + FP) divided by the total number of a dataset (examples in the dataset).

Also $\text{error_rate} = 1 - \text{accuracy}$

```
In [29]: total_misclassified = cs[0,1] + cs[1,0]
print(total_misclassified)
total_examples = cs[0,0]+cs[0,1]+cs[1,0]+cs[1,1]
print(total_examples)
print("Error rate",total_misclassified/total_examples)
print("Error rate ",1-metrics.accuracy_score(y_test,y_pred))
```

```
65
231
Error rate 0.2813852813852814
Error rate  0.2813852813852814
```

```
In [13]: print("Precision score",metrics.precision_score(y_test,y_pred))
```

```
Precision score 0.6056338028169014
```

```
In [14]: print("Recall score ",metrics.recall_score(y_test,y_pred))
```

```
Recall score  0.5375
```

```
In [15]: print("Classification report ",metrics.classification_report(y_test,y_pred))
```

Classification report		precision	recall	f1-score	support
0	0.77	0.81	0.79		151
1	0.61	0.54	0.57		80
accuracy			0.72		231
macro avg	0.69	0.68	0.68		231
weighted avg	0.71	0.72	0.71		231

Implement K-Means clustering/ hierarchical clustering on sales_data_sample.csv dataset. Determine the number of clusters using the elbow method.

```
In [198... import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
#Importing the required libraries.
```

```
In [199... from sklearn.cluster import KMeans, k_means #For clustering
from sklearn.decomposition import PCA #Linear Dimensionality reduction.
```

```
In [200... df = pd.read_csv("sales_data_sample.csv") #Loading the dataset.
```

Preprocessing

```
In [201... df.head()
```

```
Out[201]:
```

	ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	SALES	ORDERDATE
0	10107	30	95.70	2	2871.00	2/24/2003 0:00
1	10121	34	81.35	5	2765.90	5/7/2003 0:00
2	10134	41	94.74	2	3884.34	7/1/2003 0:00
3	10145	45	83.26	6	3746.70	8/25/2003 0:00
4	10159	49	100.00	14	5205.27	10/10/2003 0:00

5 rows × 25 columns

```
In [202... df.shape
```

```
Out[202]: (2823, 25)
```

```
In [203... df.describe()
```

Out[203]:

	ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	SALES	
count	2823.000000	2823.000000	2823.000000	2823.000000	2823.000000	2823.000000
mean	10258.725115	35.092809	83.658544	6.466171	3553.889072	
std	92.085478	9.741443	20.174277	4.225841	1841.865106	
min	10100.000000	6.000000	26.880000	1.000000	482.130000	
25%	10180.000000	27.000000	68.860000	3.000000	2203.430000	
50%	10262.000000	35.000000	95.700000	6.000000	3184.800000	
75%	10333.500000	43.000000	100.000000	9.000000	4508.000000	
max	10425.000000	97.000000	100.000000	18.000000	14082.800000	

In [204...

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2823 entries, 0 to 2822
Data columns (total 25 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ORDERNUMBER           2823 non-null  int64
1   QUANTITYORDERED       2823 non-null  int64
2   PRICEEACH             2823 non-null  float64
3   ORDERLINENUMBER       2823 non-null  int64
4   SALES                 2823 non-null  float64
5   ORDERDATE             2823 non-null  object
6   STATUS               2823 non-null  object
7   QTR_ID               2823 non-null  int64
8   MONTH_ID             2823 non-null  int64
9   YEAR_ID              2823 non-null  int64
10  PRODUCTLINE           2823 non-null  object
11  MSRP                  2823 non-null  int64
12  PRODUCTCODE           2823 non-null  object
13  CUSTOMERNAME          2823 non-null  object
14  PHONE                 2823 non-null  object
15  ADDRESSLINE1          2823 non-null  object
16  ADDRESSLINE2          302 non-null   object
17  CITY                  2823 non-null  object
18  STATE                 1337 non-null  object
19  POSTALCODE            2747 non-null  object
20  COUNTRY               2823 non-null  object
21  TERRITORY             1749 non-null  object
22  CONTACTLASTNAME       2823 non-null  object
23  CONTACTFIRSTNAME      2823 non-null  object
24  DEALSIZE              2823 non-null  object
dtypes: float64(2), int64(7), object(16)
memory usage: 551.5+ KB
```

In [205...

```
df.isnull().sum()
```

```
Out[205]: ORDERNUMBER      0
          QUANTITYORDERED  0
          PRICEEACH        0
          ORDERLINENUMBER  0
          SALES             0
          ORDERDATE        0
          STATUS           0
          QTR_ID           0
          MONTH_ID         0
          YEAR_ID          0
          PRODUCTLINE      0
          MSRP             0
          PRODUCTCODE      0
          CUSTOMERNAME     0
          PHONE            0
          ADDRESSLINE1     0
          ADDRESSLINE2     2521
          CITY             0
          STATE            1486
          POSTALCODE       76
          COUNTRY          0
          TERRITORY        1074
          CONTACTLASTNAME  0
          CONTACTFIRSTNAME 0
          DEALSIZE         0
          dtype: int64
```

```
In [206... df.dtypes
```

```
Out[206]: ORDERNUMBER      int64
          QUANTITYORDERED  int64
          PRICEEACH        float64
          ORDERLINENUMBER  int64
          SALES             float64
          ORDERDATE        object
          STATUS           object
          QTR_ID           int64
          MONTH_ID         int64
          YEAR_ID          int64
          PRODUCTLINE      object
          MSRP             int64
          PRODUCTCODE      object
          CUSTOMERNAME     object
          PHONE            object
          ADDRESSLINE1     object
          ADDRESSLINE2     object
          CITY             object
          STATE            object
          POSTALCODE       object
          COUNTRY          object
          TERRITORY        object
          CONTACTLASTNAME  object
          CONTACTFIRSTNAME object
          DEALSIZE         object
          dtype: object
```

```
In [207... df_drop = ['ADDRESSLINE1', 'ADDRESSLINE2', 'STATUS', 'POSTALCODE', 'CITY', 'TERRITORY']
df = df.drop(df_drop, axis=1) #Dropping the categorical unnecessary columns along with
```

```
In [208... df.isnull().sum()
```

```
Out[208]: QUANTITYORDERED    0
          PRICEEACH          0
          ORDERLINENUMBER    0
          SALES               0
          ORDERDATE          0
          QTR_ID             0
          MONTH_ID           0
          YEAR_ID            0
          PRODUCTLINE        0
          MSRP               0
          PRODUCTCODE        0
          COUNTRY            0
          DEALSIZE           0
          dtype: int64
```

```
In [209... df.dtypes
```

```
Out[209]: QUANTITYORDERED    int64
          PRICEEACH          float64
          ORDERLINENUMBER    int64
          SALES               float64
          ORDERDATE          object
          QTR_ID             int64
          MONTH_ID           int64
          YEAR_ID            int64
          PRODUCTLINE        object
          MSRP               int64
          PRODUCTCODE        object
          COUNTRY            object
          DEALSIZE           object
          dtype: object
```

```
In [ ]: # Checking the categorical columns.
```

```
In [210... df['COUNTRY'].unique()
```

```
Out[210]: array(['USA', 'France', 'Norway', 'Australia', 'Finland', 'Austria', 'UK',
                  'Spain', 'Sweden', 'Singapore', 'Canada', 'Japan', 'Italy',
                  'Denmark', 'Belgium', 'Philippines', 'Germany', 'Switzerland',
                  'Ireland'], dtype=object)
```

```
In [211... df['PRODUCTLINE'].unique()
```

```
Out[211]: array(['Motorcycles', 'Classic Cars', 'Trucks and Buses', 'Vintage Cars',
                  'Planes', 'Ships', 'Trains'], dtype=object)
```

```
In [212... df['DEALSIZE'].unique()
```

```
Out[212]: array(['Small', 'Medium', 'Large'], dtype=object)
```

```
In [213... productline = pd.get_dummies(df['PRODUCTLINE']) #Converting the categorical columns
          Dealsize = pd.get_dummies(df['DEALSIZE'])
```

```
In [214... df = pd.concat([df,productline,Dealsize], axis = 1)
```

```
In [215... df_drop = ['COUNTRY', 'PRODUCTLINE', 'DEALSIZE'] #Dropping Country too as there are
          df = df.drop(df_drop, axis=1)
```

```
In [216... df['PRODUCTCODE'] = pd.Categorical(df['PRODUCTCODE']).codes #Converting the datatype
```



```
In [217...] df.drop('ORDERDATE', axis=1, inplace=True) #Dropping the Orderdate as Month is already present
```


```
In [218...] df.dtypes #ALL the datatypes are converted into numeric
```

```
Out[218]: QUANTITYORDERED    int64
PRICEEACH                float64
ORDERLINENUMBER          int64
SALES                    float64
QTR_ID                   int64
MONTH_ID                 int64
YEAR_ID                 int64
MSRP                     int64
PRODUCTCODE              int8
Classic Cars              uint8
Motorcycles               uint8
Planes                   uint8
Ships                    uint8
Trains                   uint8
Trucks and Buses         uint8
Vintage Cars              uint8
Large                    uint8
Medium                   uint8
Small                    uint8
dtype: object
```

Plotting the Elbow Plot to determine the number of clusters.

```
In [219...] distortions = [] # Within Cluster Sum of Squares from the centroid
K = range(1,10)
for k in K:
    kmeanModel = KMeans(n_clusters=k)
    kmeanModel.fit(df)
    distortions.append(kmeanModel.inertia_) #Appending the inertia to the Distortions list
```

```
In [220...] plt.figure(figsize=(16,8))
plt.plot(K, distortions, 'bx-')
plt.xlabel('k')
plt.ylabel('Distortion')
plt.title('The Elbow Method showing the optimal k')
plt.show()
```



As the number of k increases Inertia decreases.

Observations: A Elbow can be observed at 3 and after that the curve decreases gradually.

```
In [221...] X_train = df.values #Returns a numpy array.
```

```
In [222...] X_train.shape
```

```
Out[222]: (2823, 19)
```

```
In [223...] model = KMeans(n_clusters=3,random_state=2) #Number of cluster = 3
model = model.fit(X_train) #Fitting the values to create a model.
predictions = model.predict(X_train) #Predicting the cluster values (0,1,or 2)
```

```
In [225...] unique,counts = np.unique(predictions,return_counts=True)
```

```
In [226...] counts = counts.reshape(1,3)
```

```
In [227...] counts_df = pd.DataFrame(counts,columns=['Cluster1','Cluster2','Cluster3'])
```

```
In [228...] counts_df.head()
```

```
Out[228]:
```

	Cluster1	Cluster2	Cluster3
0	1083	1367	373

Visualization

```
In [229...] pca = PCA(n_components=2) #Converting all the features into 2 columns to make it ec
```

```
In [230...] reduced_X = pd.DataFrame(pca.fit_transform(X_train),columns=['PCA1','PCA2']) #Creat
```

```
In [231...] reduced_X.head()
```


Out[231]:

	PCA1	PCA2
0	-682.488323	-42.819535
1	-787.665502	-41.694991
2	330.732170	-26.481208
3	193.040232	-26.285766
4	1651.532874	-6.891196

In [232...

```
#Plotting the normal Scatter Plot  
plt.figure(figsize=(14,10))  
plt.scatter(reduced_X['PCA1'],reduced_X['PCA2'])
```

Out[232]: <matplotlib.collections.PathCollection at 0x218dc747880>



In [233...

```
model.cluster_centers_ #Finding the centriods. (3 Centriods in total. Each Array co
```

```
Out[233]: array([[ 3.72031394e+01,  9.52120960e+01,  6.44967682e+00,
  4.13868425e+03,  2.72022161e+00,  7.09879963e+00,
  2.00379409e+03,  1.13248384e+02,  5.04469067e+01,
  3.74884580e-01,  1.15420129e-01,  9.41828255e-02,
  8.21791320e-02,  1.84672207e-02,  1.16343490e-01,
  1.98522622e-01,  2.08166817e-17,  1.00000000e+00,
 -6.66133815e-16],
 [ 3.08302853e+01,  7.00755230e+01,  6.67300658e+00,
  2.12409474e+03,  2.71762985e+00,  7.09509876e+00,
  2.00381127e+03,  7.84784199e+01,  6.24871982e+01,
  2.64813460e-01,  1.21433797e-01,  1.29480614e-01,
  1.00219459e-01,  3.87710315e-02,  9.21726408e-02,
  2.53108998e-01,  6.93889390e-18,  6.21799561e-02,
  9.37820044e-01],
 [ 4.45871314e+01,  9.98931099e+01,  5.75603217e+00,
  7.09596863e+03,  2.71045576e+00,  7.06434316e+00,
  2.00389008e+03,  1.45823056e+02,  3.14959786e+01,
  5.33512064e-01,  1.07238606e-01,  7.23860590e-02,
  2.14477212e-02,  1.07238606e-02,  1.31367292e-01,
  1.23324397e-01,  4.20911528e-01,  5.79088472e-01,
  5.55111512e-17]])
```


```
In [234... reduced_centers = pca.transform(model.cluster_centers_) #Transforming the centroids
```

```
In [235... reduced_centers
```

```
Out[235]: array([[ 5.84994044e+02, -4.36786931e+00],
 [-1.43005891e+03,  2.60041009e+00],
 [ 3.54247180e+03,  3.15185487e+00]])
```

```
In [236... plt.figure(figsize=(14,10))
plt.scatter(reduced_X['PCA1'],reduced_X['PCA2'])
plt.scatter(reduced_centers[:,0],reduced_centers[:,1],color='black',marker='x',s=36)
```

```
Out[236]: <matplotlib.collections.PathCollection at 0x218deb6e220>
```



In [237... `reduced_X['Clusters'] = predictions` *#Adding the Clusters to the reduced dataframe.*

In [238... `reduced_X.head()`

Out[238]:


	PCA1	PCA2	Clusters
0	-682.488323	-42.819535	1
1	-787.665502	-41.694991	1
2	330.732170	-26.481208	0
3	193.040232	-26.285766	0
4	1651.532874	-6.891196	0

In [239... *#Plotting the clusters*

```
plt.figure(figsize=(14,10))
#               taking the cluster number and first column               taking the cluster number and first column
plt.scatter(reduced_X[reduced_X['Clusters'] == 0].loc[:, 'PCA1'], reduced_X[reduced_X['Clusters'] == 0].loc[:, 'PCA2'], color='red', marker='x', s=300)
plt.scatter(reduced_X[reduced_X['Clusters'] == 1].loc[:, 'PCA1'], reduced_X[reduced_X['Clusters'] == 1].loc[:, 'PCA2'], color='blue', marker='x', s=300)
plt.scatter(reduced_X[reduced_X['Clusters'] == 2].loc[:, 'PCA1'], reduced_X[reduced_X['Clusters'] == 2].loc[:, 'PCA2'], color='green', marker='x', s=300)

plt.scatter(reduced_centers[:,0], reduced_centers[:,1], color='black', marker='x', s=300)
```

Out[239]: <matplotlib.collections.PathCollection at 0x218dce9e1f0>



In []: